

Inf2C - Computer Systems

Lecture 2

Data Representation

Boris Grot

School of Informatics
University of Edinburgh



Last lecture

- Course overview
 - Piazza: up & running. Use it!
 - Tutorials: start in week 3
 - Labs: drop-in. Start in week 3
- Moore's law
- Types of computer systems
- Computer components
- Computer system stack



Lecture 2: Data Representation

- The way in which data is represented in computer hardware affects
 - complexity of circuits
 - cost
 - speed
 - reliability
- Must consider how to design hardware for
 - Storing data: memory
 - Manipulating data: processing (e.g., adders, multipliers)

Lecture outline

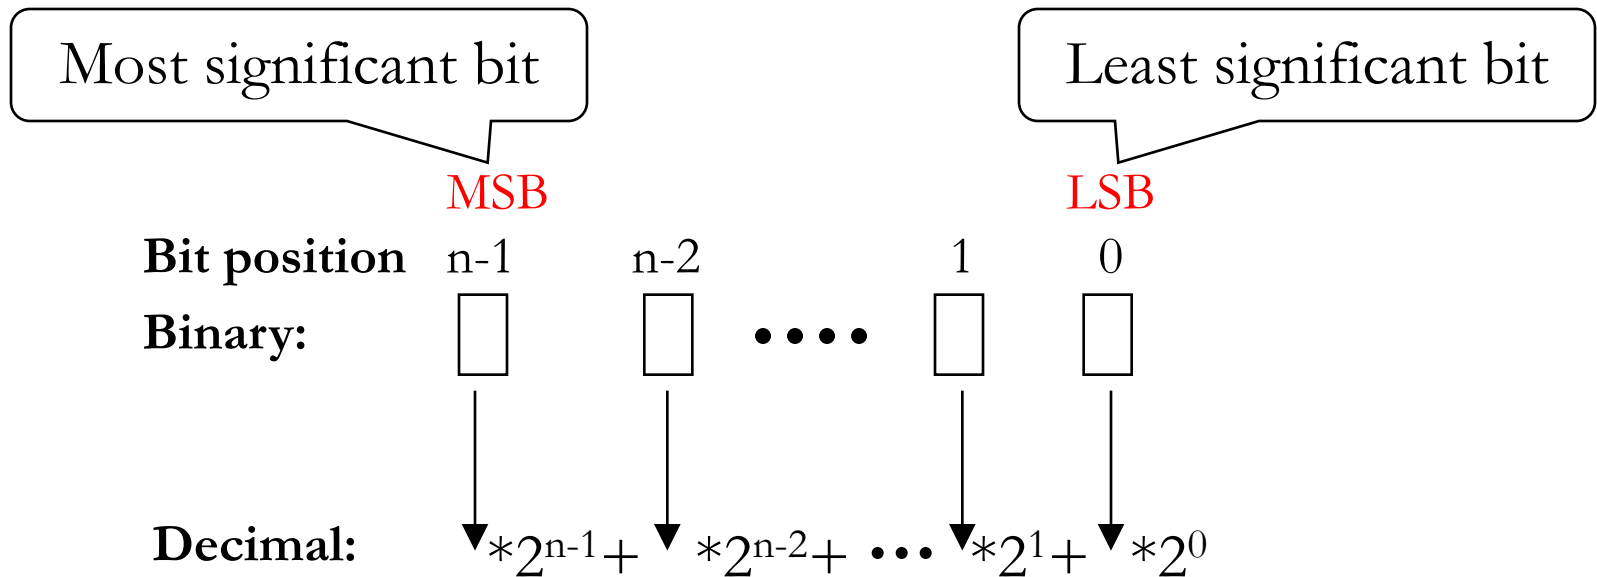
- The bit – atomic unit of data
- Representing numbers
 - Integers
 - Floating point
- Representing text

The bit

- Information represented as sequences of symbols
 - Humans use letters, numerals, punctuation, whitespace
 - Computers use just 0s and 1s, *bits*
- *Bit* – an acronym for Binary digiT
- Advantages: easy to do computation, very reliable, simple & reusable circuits
- Disadvantages: little information per bit → must use many bits. $512 \equiv 1\ 0000\ 0000$, 'A' $\equiv 0100\ 0001$

Natural numbers representation

- Non-negative (unsigned) integers are very simple to represent in binary



Basic operations

- Addition, subtraction with unsigned binary numbers is easy:

$$\begin{array}{r} \textcolor{red}{1111} \\ 01101 \\ +01011 \\ \hline 11000 \end{array} \quad \begin{array}{r} \textcolor{red}{0010} \\ 01101 \\ -01011 \\ \hline 00010 \end{array}$$

Diagram illustrating binary addition and subtraction with bit counts:

- The addition of 01101 and 01011 results in 11000 . The number of bits in the operands is 13, and the result is 24 bits long.
- The subtraction of 01011 from 01101 results in 00010 . The number of bits in the operands is 11, and the result is 2 bits long.

Fixed bit-length arithmetic

- Hardware cannot handle infinitely long bit sequences
- We end up with a few fixed-size data types
 - **Byte**: always 8 bits
 - **Word**: the typical unit of data on which a processor operates (32 or 64 bits most common today)
- **Overflow** happens when a result does not fit
 - Numbers wrap-around when they become too large
 - Arithmetic is modulo 2^N , where N =number of bits



What about negative numbers?

- **Sign-magnitude** representation:
 - Use 1st bit (MSB) as the sign
 - 1 → negative, 0 → positive
 - $0010 \equiv 2$ $1010 \equiv -2$
- Complicates addition and subtraction
 - The actual operation depends on the sign
- Has positive and negative zero
 - $0000 \equiv 0$ $1000 \equiv -0$



Better way: **2's complement** representation

Two's complement: the intuition

- Want: $X + (-X) = 0$
- Insight: don't need the full sum to be 0
 - Only the bits that can be represented within a computer's fixed width need to be 0
- Approach:
 - Represent the negation of X as $2^N - X$
 - Then: $X + (-X) = X + (2^N - X) = 2^N$
 - Recall: largest number represented with N bits: $2^N - 1$
 - Note that N lowest bits of the sum are all 0!



Two's complement: example

Given:

- 3-bit fixed width ($N=3$)
- $X = 2$ (decimal) \rightarrow 0 1 0 (binary)

$$2^N = 8 \text{ (dec)} \rightarrow 1 \ 0 \ 0 \ 0 \text{ (bin)}$$

$$-X = 2^N - X = 8 - 2 = 6 \text{ (dec)} \rightarrow 1 \ 1 \ 0 \text{ (bin)}$$

Check:

$$X + (-X) = 0 \ 1 \ 0 + 1 \ 1 \ 0 = 1 \ 0 \ 0 \ 0$$

3-bit fixed width



Efficiently computing 2's complement

EASY!

“Flip the bits and add 1”

Example:

$$X = 0\ 1\ 0\ (\text{bin}) \rightarrow 2\ (\text{dec})$$

Flip the bits: 1 0 1


$$\text{Add 1:} \quad 1\ 1\ 0\ (\text{bin}) \rightarrow -X$$



The roots of the idea

John von Neumann (died: 1957)

- Co-inventor of the stored program concept
- Proposed 2's complement idea in a 1945 paper
- Also came up with cellular automata, numerical weather forecasting, concept of global warming
- Outside of computing: linear programming, quantum logic, policy of mutually assured destruction, and more!



**Patron saint of
this class**

2's complement details

- The MSB is the sign
- $A - B = A + (2\text{'s complement of } B)$
- Arithmetic operations do not depend on the operands' signs
- Range is asymmetric: -2^{n-1} to $2^{n-1}-1$
- There are two kinds of overflows:
 - Positive overflow produces a negative number
 - Negative **underflow** produces a positive number



Converting between data types

- Converting a 2's complement number from a smaller to a larger representation is done by **sign extension**

Example: from byte to short (16 bits):

$2 = 00000010 \Rightarrow ? ? ? ? ? ? ? ? 00000010$

$-2 = 11111110 \Rightarrow ? ? ? ? ? ? ? ? 11111110$

$2 = \overbrace{00000010}^{\text{(byte)}} \Rightarrow \overbrace{00000000}^{\text{(byte)}} \overbrace{00000010}^{\text{(byte)}}$
(byte) (short)

$-2 = \overbrace{11111110}^{\text{(byte)}} \Rightarrow \overbrace{11111111}^{\text{(byte)}} \overbrace{11111110}^{\text{(byte)}}$
(byte) (short)

Shifting

- Shifting: move the bits of a data type left or right
 - Data bits falling off the edge are lost
- For left shifts, 0s fill in the empty bit places
- For right shifts, two options:
 - Fill with 0 (**logical shift**): for non-numerical data
 - Fill with MSB (**arithmetic shift**): for 2's complement numbers
- Shift left by n is equivalent to multiplying by 2^n
- Shift right by n is equivalent to dividing by 2^n and rounding towards $-\infty$
- Example
 - $6 = 00000110 \gg 2 \rightarrow 00000001 = 1$
 - $-6 = 11111010 \gg 2 \rightarrow 11111110 = -2$



Hexadecimal notation

- Binary numbers (and other binary-encoded information) are too long and tedious for us to use
- Solution: use a more compact encoding
 - Hexadecimal (base 16) is most common
- Hex digits: 0-9 and A-F
 - $A=10_{\text{dec}}$, $B=11$, ..., $F=15$
- Conversion to/from binary is very easy:
Every 4 bits correspond to 1 hex digit:

$$\begin{array}{ccccccc} 1 & 1 & 1 & 1 & 1 & 0 & 0 & 0 \\ \underbrace{\hspace{1.5cm}} & \underbrace{\hspace{1.5cm}} & & & & & & \\ F(15) & 8 & & & & & & \end{array} = 0xF8$$

Hex is just a convenience for humans
Computers use the binary form

Real numbers - floating point

- Java's **float** (32 bits)
double (64 bits)

- Binary representation:

– example 0.75 in base 10 \Rightarrow 0.11 in base 2

$$\begin{array}{c} \swarrow \quad \searrow \\ (2^{-1} + 2^{-2} = 0.5 + 0.25 = 0.75) \end{array}$$

- example **0.75** in base 10 \Rightarrow **0.11** in base 2

■ Normalization:

$0.11 \Rightarrow 1.1 \times 2^{-1}$

Mantissa
(aka significand)

exponent

implicit
(always 1)



Why normalize?

Three reasons:

1. Simplifies machine representation
(don't need to represent the fraction separator)
2. Simplifies comparisons
 - Which one is bigger: 0.0000101 or 0.000001 ?
 - 1.01×2^{-5} vs 1.0×2^{-6}
3. Is more compact for very small/large numbers
 - E.g., $0.000000000000000001 = 1.0 \times 2^{-16}$
 - or can be made more compact (by rounding fraction)

Floating point conversion example #1

Convert the number 25 to floating point with normalization

1) 25 in base 10 \Rightarrow 11001 in base 2

2) 11001 to normalized floating point $\Rightarrow 1.1001 \times 2^4$


Understand that:

- The number is normalized
- 1.1001 is mantissa (aka **significand**)
- 4 is exponent
- sign is “+” (implicit here)



IEEE 754 Floating Point standard

- Need a standard to represent and compute with fixed-width floats
- 32 bit representation:



$$(-1)^s \times (1.m) \times 2^{\text{exp}-127} \quad \text{Bias}$$

e.g.,

$$(0.75)_{10} \rightarrow (0.11)_2 \rightarrow (1.1 \times 2^{-1})_2$$


$$\rightarrow s = 0, m = 1, \text{exp} = 126$$

$$\rightarrow 0 \ 01111110 \ 100000000000000000000000$$

Note: representation does NOT use 2's complement

IEEE 754 Floating Point standard

- Need a standard to represent and compute with fixed-width floats
- 32 bit representation:



$$(-1)^s \times (1.m) \times 2^{\text{exp}-127} \quad \text{Bias}$$

e.g.,

$$(0.75)_{10} \rightarrow (0.11)_2 \rightarrow (1.1 \times 2^{-1})_2$$

$$\rightarrow s = 0, m = 1, \text{exp} = 126$$

$$\rightarrow 0 \ 01111110 \ 100000000000000000000000$$

- 64 bit representation:
 - exponent = 11 bits; mantissa = 52 bits



IEEE 754 Floating Point standard

- Why bias?
 - Avoids the complexity of $+/-$ exponents
 - Simplifies relative ordering of FP numbers
- Note: processors usually have specialized floating point units to perform FP arithmetic

IEEE 754 floating point conversion #2

Example: Convert 23.5 (decimal) to IEEE 754 floating point

Start: 23 in base 10 \Rightarrow 10111 in base 2

IEEE 754 floating point conversion #2


Example: Convert 23.5 (decimal) to IEEE 754 floating point

Start: 23 in base 10 \Rightarrow 10111 in base 2

- 1) 23.5 in base 10 \Rightarrow 10111.1 in base 2
- 2) 10111.1 to normalized floating point $\Rightarrow 1.01111 \times 2^4$
- 3) S = 0

M = 01111 is mantissa (remember: 1. is implicit)

Exp = $4+127 = 131$ in base 10 \Rightarrow 1000 0011 in base 2



IEEE 754 Floating Point notation

Exponent	Mantissa	Meaning
0	0	0
1-254	Anything	Floating point number
255	0	Infinity (signed)
255	Non-zero	Not-a-number (NaN)

32-bit representation

Representing characters

- Characters need to be encoded in binary too
- Operations on characters have simpler requirements than on numbers, so the encoding choice is not crucial
- Most common representation is ASCII
 - Each character is held in a byte
 - E.g. '0' is 0x30, 'A' is 0x41, 'a' is 0x61
- Java uses Unicode which can encode characters from many (all?) languages
 - 16 bits per character required

Representing strings

- Words, sentences, etc. are just **strings** of characters
- How is the end of a string identified?
 - No common standard exists. Different programming languages use different encodings
 - In C: a special character, encoded as 0x00
 - Also called NULL character
 - In Java: string length is kept with the string itself
 - string is an object and length is one of its member variables

Summary

- Computers use binary representation
- Signed numbers: sign-magnitude vs 2's complement
- Floating point
- Characters and strings